

```
In [ ]: from keras.preprocessing.image import ImageDataGenerator
import numpy as np
import pandas as pd
import tensorflow as tf
from sklearn.model_selection import train_test_split
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense, InputLayer, Conv2D, MaxPooling2D
from tensorflow.keras.models import Model
from keras.applications.vgg16 import VGG16
from tensorflow.keras.callbacks import EarlyStopping, Callback
```

```
In [ ]: from google.colab import files
uploaded = files.upload()
for fn in uploaded.keys():
    print('uploaded file "{(name)" with length {length} bytes'.format(
        name=fn, length=len(uploaded[fn])))
```

```
Choose Files | No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving labels_final.csv to labels_final (1).csv
uploaded file "labels_final.csv" with length 2188171 bytes
```

```
In [ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
In [ ]: !pip install pyunpack
!pip install patool
from pyunpack import Archive
Archive('/content/drive/MyDrive/rv1-cdip.rar').extractall('/content')

Collecting pyunpack
  Downloading pyunpack-0.2.2-py2.py3-none-any.whl (3.8 kB)
Collecting entrypoint2
  Downloading entrypoint2-0.2.4-py3-none-any.whl (6.2 kB)
Collecting easyprocess
  Downloading EasyProcess-0.3-py2.py3-none-any.whl (7.9 kB)
Installing collected packages: entrypoint2, easyprocess, pyunpack
Successfully installed easyprocess-0.3 entrypoint2-0.2.4 pyunpack-0.2.2
Collecting patool
  Downloading patool-1.12-py2.py3-none-any.whl (77 kB)
    77 kB 3.2 MB/s
Installing collected packages: patool
Successfully installed patool-1.12
```

```
In [ ]: df=pd.read_csv("labels_final.csv",dtype=str)
```

```
In [ ]: df.head(5)
```

	path	label
0	images/vlv/hvh71d00509132755+2755.tif	3
1	images/l/vb/hd19d00502213303.tif	3
2	images/x/el/dx/d05a002075325674.tif	2
3	images/oj/b/bq/b60d00517511301+1301.tif	3
4	images/q/z/k/qz/k17e002031320195.tif	7

```
In [ ]: datagen=ImageDataGenerator(rescale=1./255.,validation_split=0.25)
```

```
In [ ]: train_generator = datagen.flow_from_dataframe(
    dataframe=df,
    directory="/content/data_final",
    x_col="path",
    y_col="label",
    has_ext=False,
    subset="training",
    batch_size=32,
    shuffle=True,
    class_mode="categorical",
    target_size=(224, 224))

validation_generator=datagen.flow_from_dataframe(
    dataframe=df,
    directory="/content/data_final",
    x_col="path",
    y_col="label",
    subset="validation",
    batch_size=32,
    seed=42,
    shuffle=True,
    class_mode="categorical",
    target_size=(224, 224))

Found 36000 validated image filenames belonging to 16 classes.
Found 12000 validated image filenames belonging to 16 classes.
```

```
In [ ]: import datetime
import os
! rm -rf ./logs/
logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
print(logdir)
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir,histogram_freq=1, write_graph=True)
```

logs/20210928-182507  
Model 1 Training

```
In [ ]: vgg=VGG16(weights='imagenet', include_top=False,input_shape = (224, 224, 3))
# don't train existing weights
for layer in vgg.layers:
    layer.trainable=False

top_model = Conv2D(filters=32, kernel_size=(3,3), activation='relu',padding='same',kernel_initializer=tf.keras.initializers.HeUniform()))(vgg.output)
top_model=MaxPooling2D(pool_size=(2,2),strides=2)(top_model)
# Flatten the output layer to 1 dimension
top_model = Flatten()(top_model)
top_model = Dense(64, activation='relu',kernel_initializer=tf.keras.initializers.HeUniform()))(top_model)
top_model = Dense(32, activation='relu',kernel_initializer=tf.keras.initializers.HeUniform()))(top_model)
output_layer = Dense(16, activation='softmax')(top_model)

model = Model(inputs = vgg.input, outputs = output_layer)
print(model.summary())
```

Exception ignored in: <function IteratorResourceDeleter.\_\_del\_\_ at 0x7f6c6e67e8c0>  
Traceback (most recent call last):  
File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/data/ops/iterator\_ops.py", line 546, in \_\_del\_\_  
handle=self.handle, deleter=self.deleter)  
File "/usr/local/lib/python3.7/dist-packages/tensorflow/python/ops/gen\_dataset\_ops.py", line 1264, in delete\_iterator  
ctx, "DeleteIterator", name, handle, deleter)  
KeyboardInterrupt:  
Model: "model\_1"

Layer (type)	Output Shape	Param #
input_8 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d_16 (Conv2D)	(None, 7, 7, 32)	147488
max_pooling2d_1 (MaxPooling2D)	(None, 3, 3, 32)	0
Flatten_0 (Flatten)	(None, 288)	0
dense_10 (Dense)	(None, 64)	18496
dense_11 (Dense)	(None, 32)	2880
dense_12 (Dense)	(None, 16)	528
=====		
Total params:	14,883,280	
Trainable params:	168,592	
Non-trainable params:	14,714,688	
None		

```
In [ ]: # Before training a model, we need to configure the learning process, which is done via the compile method
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy'])
# have to stop the training if validation accuracy is not increased in last 2 epochs.
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.20, patience=4, verbose=1)
model.fit(train_generator, validation_data = validation_generator, steps_per_epoch=163, epochs = 10, callbacks=[earlystop, tensorboard_callback])

Epoch 1/10
163/163 [=====] - 130s 792ms/step - loss: 2.5451 - accuracy: 0.2203 - val_loss: 2.2867 - val_accuracy: 0.3320
Epoch 2/10
163/163 [=====] - 128s 786ms/step - loss: 2.1003 - accuracy: 0.3815 - val_loss: 1.9501 - val_accuracy: 0.4371
Epoch 3/10
163/163 [=====] - 128s 785ms/step - loss: 1.8454 - accuracy: 0.4482 - val_loss: 1.7446 - val_accuracy: 0.4883
Epoch 4/10
163/163 [=====] - 128s 786ms/step - loss: 1.6850 - accuracy: 0.4956 - val_loss: 1.6221 - val_accuracy: 0.5196
Epoch 5/10
163/163 [=====] - 149s 915ms/step - loss: 1.5693 - accuracy: 0.5200 - val_loss: 1.5567 - val_accuracy: 0.5414
Epoch 6/10
163/163 [=====] - 147s 994ms/step - loss: 1.4928 - accuracy: 0.5546 - val_loss: 1.4863 - val_accuracy: 0.5563
Epoch 7/10
163/163 [=====] - 148s 911ms/step - loss: 1.4335 - accuracy: 0.5796 - val_loss: 1.4435 - val_accuracy: 0.5670
Epoch 8/10
163/163 [=====] - 146s 900ms/step - loss: 1.4045 - accuracy: 0.5937 - val_loss: 1.3994 - val_accuracy: 0.5846
Epoch 9/10
163/163 [=====] - 143s 879ms/step - loss: 1.3382 - accuracy: 0.6817 - val_loss: 1.3638 - val_accuracy: 0.5923
Epoch 00000: early stopping
Epoch 00000: early stopping
<keras.callbacks.History at 0x7f6c59b29510>
```

```
In [ ]: %load_ext tensorboard
%tensorboard --logdir $logdir
```

The tensorboard extension is already loaded. To reload it, use:  
%reload\_ext tensorboard

## Model-2

1. Use VGG-16 pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size 7x7x512 can be equivalently expressed as a CONV layer with F=7,P=0,S=1,K=4096. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be 1x1x4096 since only a single depth column "fits" across the input volume, giving identical result as the initial FC layer. You can refer [this](#) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top Layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**
3. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

```
In [ ]: vgg_model=VGG16(weights='imagenet', include_top=False,input_shape = (224, 224, 3))
# we are not training existing weights
model.trainable=False
top_model = Conv2D(4096,kernel_size=(1,1),padding='valid', activation='relu',kernel_initializer=tf.keras.initializers.HeUniform()))(vgg_model.output)
top_model=Conv2D(4096,kernel_size=(1,1),padding='valid', activation='relu')(top_model)
# Flatten the output layer to 1 dimension
top_model = Flatten()(top_model)
output_layer = Dense(16, activation='softmax')(top_model)
model = Model(inputs = vgg_model.input, outputs = output_layer)
print(model.summary())
```

Model: "model_2"		
Layer (type)	Output Shape	Param #
input_4 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d_5 (Conv2D)	(None, 7, 7, 4096)	2101248
conv2d_6 (Conv2D)	(None, 7, 7, 4096)	16781312
Flatten_2 (Flatten)	(None, 200704)	0
dense_2 (Dense)	(None, 16)	3211280
=====		
Total params:	36,800,528	
Trainable params:	22,093,840	
Non-trainable params:	14,714,688	
None		

```
In [ ]: # Before training a model, we need to configure the learning process, which is done via the compile method
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy'])
# have to stop the training if validation accuracy is not increased in last 2 epochs.
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.20, patience=2, verbose=1)
model.fit(train_generator, validation_data = validation_generator, steps_per_epoch=220, epochs = 5, callbacks=[earlystop, tensorboard_callback])

Epoch 1/5
220/220 [=====] - 231s 052ms/step - loss: 1.7850 - accuracy: 0.4756 - val_loss: 1.3178 - val_accuracy: 0.6050
Epoch 2/5
220/220 [=====] - 186s 047ms/step - loss: 1.2206 - accuracy: 0.6200 - val_loss: 1.2061 - val_accuracy: 0.6433
Epoch 3/5
220/220 [=====] - 186s 046ms/step - loss: 1.1177 - accuracy: 0.6595 - val_loss: 1.0847 - val_accuracy: 0.6777
Epoch 00000: early stopping
Epoch 00000: early stopping
<keras.callbacks.History at 0x7f6c50287dd0>
```

```
In [ ]: %load_ext tensorboard
%tensorboard --logdir $logdir
```

## Model-3

1. Use same network as Model-2 'INPUT --> VGG-16 without Top Layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

```
In [ ]: vgg =VGG16(weights='imagenet', include_top=False,input_shape = (224, 224, 3))
# freeze convolution blocks training only last 0 layer
for layer in vgg.layers[:13]:
    layer.trainable = False

for i, layer in enumerate(vgg.layers):
    print(i, layer.name, layer.trainable)

0 input_6 False
1 block1_conv1 False
2 block1_conv2 False
3 block1_pool False
4 block2_conv1 False
5 block2_conv2 False
6 block2_pool False
7 block3_conv1 False
8 block3_conv2 False
9 block3_conv3 False
10 block3_pool False
11 block4_conv1 False
12 block4_conv2 False
13 block4_conv3 True
14 block4_pool True
15 block5_conv1 True
16 block5_conv2 True
17 block5_conv3 True
18 block5_pool True
```

```
In [ ]: top_model = Conv2D(4096,kernel_size=(1,1),padding='valid', activation='relu',kernel_initializer=tf.keras.initializers.HeUniform()))(vgg._output)
top_model=Conv2D(4096,kernel_size=(1,1),padding='valid', activation='relu')(top_model)
# Flatten the output layer to 1 dimension
top_model = Flatten()(top_model)
output_layer = Dense(16, activation='softmax')(top_model)
model = Model(inputs = vgg.input, outputs = output_layer)
print(model.summary())

Model: "model_6"
Layer (type) Output Shape Param #
=====
input_6 (InputLayer) [(None, 224, 224, 3)] 0
block1_conv1 (Conv2D) (None, 224, 224, 64) 1792
block1_conv2 (Conv2D) (None, 224, 224, 64) 36928
block1_pool (MaxPooling2D) (None, 112, 112, 64) 0
block2_conv1 (Conv2D) (None, 112, 112, 128) 73856
block2_conv2 (Conv2D) (None, 112, 112, 128) 147584
block2_pool (MaxPooling2D) (None, 56, 56, 128) 0
block3_conv1 (Conv2D) (None, 56, 56, 256) 295168
block3_conv2 (Conv2D) (None, 56, 56, 256) 590080
block3_conv3 (Conv2D) (None, 56, 56, 256) 590080
block3_pool (MaxPooling2D) (None, 28, 28, 256) 0
block4_conv1 (Conv2D) (None, 28, 28, 512) 1180160
block4_conv2 (Conv2D) (None, 28, 28, 512) 2359808
block4_conv3 (Conv2D) (None, 28, 28, 512) 2359808
block4_pool (MaxPooling2D) (None, 14, 14, 512) 0
block5_conv1 (Conv2D) (None, 14, 14, 512) 2359808
block5_conv2 (Conv2D) (None, 14, 14, 512) 2359808
block5_conv3 (Conv2D) (None, 14, 14, 512) 2359808
block5_pool (MaxPooling2D) (None, 7, 7, 512) 0
conv2d_13 (Conv2D) (None, 7, 7, 4096) 2101248
conv2d_14 (Conv2D) (None, 7, 7, 4096) 16781312
flatten_6 (Flatten) (None, 200704) 0
dense_6 (Dense) (None, 16) 3211280
=====
Total params: 36,800,528
Trainable params: 31,533,072
Non-trainable params: 5,275,456
None
```

```
In [ ]: # Before training a model, we need to configure the learning process, which is done via the compile method
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.0001),
    loss='categorical_crossentropy',
    metrics=['accuracy'])
# have to stop the training if validation accuracy is not increased in last 2 epochs.
earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0.20, patience=2, verbose=1)
model.fit(train_generator, validation_data = validation_generator, steps_per_epoch=163, epochs = 10, callbacks=[earlystop, tensorboard_callback])

Epoch 1/10
163/163 [=====] - 177s 1s/step - loss: 2.1504 - accuracy: 0.3317 - val_loss: 1.7481 - val_accuracy: 0.4532
Epoch 2/10
163/163 [=====] - 175s 1s/step - loss: 1.5033 - accuracy: 0.5303 - val_loss: 1.3989 - val_accuracy: 0.5672
Epoch 3/10
163/163 [=====] - 175s 1s/step - loss: 1.3339 - accuracy: 0.5964 - val_loss: 1.3170 - val_accuracy: 0.6037
Epoch 00000: early stopping
Epoch 00000: early stopping
<keras.callbacks.History at 0x7f6b6128bed0>
```

```
In [ ]: %load_ext tensorboard
%tensorboard --logdir $logdir
```

The tensorboard extension is already loaded. To reload it, use:  
%reload\_ext tensorboard

TensorBoard Observation

Model 1

a. We can see from the loss function at every epoch. Since the loss drops very quickly and stabilizes after the 2nd epochs, it indicates that our model has learned

b. We can see from the Accuracy function at every epoch. accuracy is not much improving .

Model 2

a. We can see from the loss function at every epoch. Since the loss drops very quickly and stabilizes after the 2nd epochs, it indicates that our model has learned

b. We can see from the Accuracy function at every epoch. accuracy is not much improving .

Model 3

a. We can see from the loss function at every epoch. Since the loss drops very quickly and stabilizes after the 2nd epochs, it indicates that our model has learned

b. We can see from the Accuracy function at every epoch. accuracy is not much improving .